

Shenandoah GC 2.0

Roman Kennke
@rkennke

Overview

- Shenandoah introduction
- Old vs new barrier scheme
- New platforms
- Elimination of forwarding pointer word
- Self-fixing barriers

Shenandoah Introduction

Shenandoah is:

- A garbage collector in OpenJDK
- Concurrent GC
- Universal
- Aimed to solve the GC pause problem
 - Regardless of heap size
 - Slow/tiny hardware, cloud, constrained environments, etc

Shenandoah Introduction

You may have heard:

- That Shenandoah requires more memory
- Has complicated barriers
- Lacks some optimizations

→ No longer true

Shenandoah Introduction

Young

Old

Ser/Par

Copy

Mark

Compact

CMS

Copy

Conc-Mark

Conc-Sweep

G1

Copy

Conc-Mark

Compact

Shenandoah
ZGC


Conc-Mark

Conc-Compact

Concurrent Marking

Not covered here, no time, sorry

If you're interested, it's using "Snapshot at the beginning" algorithm, similar to G1



Concurrent Compaction


Compaction:

1. Copy all reachable objects from the collection set to empty regions
2. Adjust all references to old copies to point to new copies
→ This is easy while the program is stopped.

Problem: how to do this while program is running?

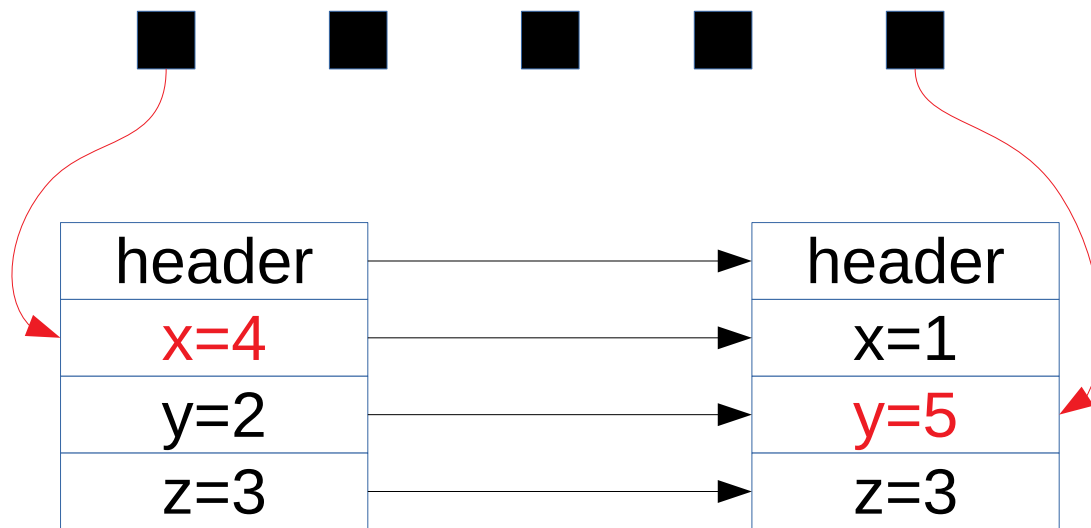


Concurrent Compaction




Concurrent Compaction

Problem:



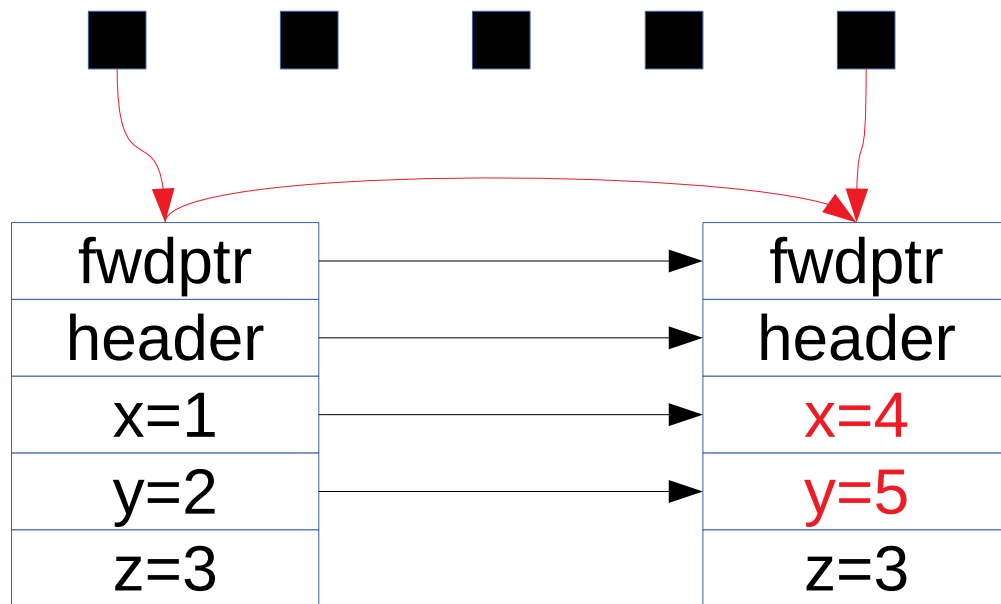
Concurrent Compaction

Solution:



Concurrent Compaction

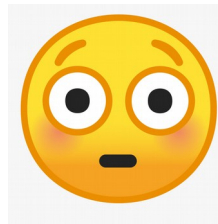
Solution (Shenandoah 1.0):



Shenandoah Barriers

Requires:

- Read barriers: Intercept reads, resolve by reading forwarding pointer.
- Write barriers: Intercept writes, establish canonical copy by CASing forwarding pointer
- Object-Equals barriers: to ensure correct result when comparing two copies of same object
- CAS barriers: combined problem of above
- Arraycopy barriers: Special case for bulk array copy



Shenandoah Barriers

Example:

```
void doSomething() {  
    Foo foo = bar.foo;  
    while (...) { // hot loop  
        a = foo.x;  
        foo.y = b;  
    }  
}
```

Shenandoah Barriers

Example:

```
void doSomething() {  
    Foo foo = bar.foo;  
    while (...) { // hot loop  
        foo' = rb(foo); // Required for primitive fields too  
        a = foo'.x;  
        foo'' = wb(foo); // Required for primitive fields too  
        foo''.y = b;  
    }  
}
```

Shenandoah Barriers

We can optimize this with some JIT heroics:

```
void doSomething() {  
    Foo foo = bar.foo;  
    foo' = rb(foo);  
    foo'' = wb(foo);  
    while (...) { // hot loop  
        a = foo'.x;  
        foo''.y = b;  
    }  
}
```

Shenandoah Barriers

We can optimize this with some JIT heroics:

```
void doSomething() {  
    Foo foo = bar.foo;  
    foo' = wb(foo);  
    while (...) { // hot loop  
        a = foo'.x;  
        foo'.y = b;  
    }  
}
```


Shenandoah Barriers

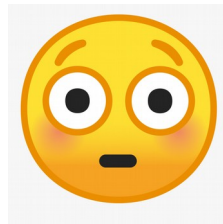
Idea: Establish canonical copy right when object is loaded

```
void doSomething() {  
    Foo foo' = bar.foo;  
    foo= lrb(foo');  
    while (...) { // hot loop  
        a = foo.x;  
        foo.y = b;  
    }  
}
```

Shenandoah 2.0 Barriers: LRB

Load-reference-barriers:

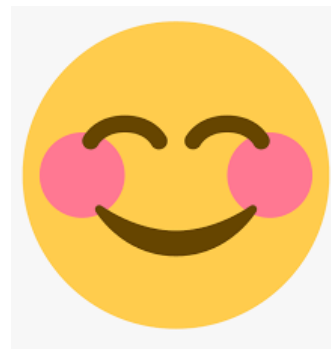
- ~~Read-barriers: Intercept reads, resolve by reading forwarding pointer.~~
- Write-barriers: Intercept ~~writes~~ **reference-loads**, establish canonical copy by CASing forwarding pointer
- ~~Above two also required for primitive fields~~
- ~~Object-equals-barriers: to ensure correct result when comparing two copies of same object~~
- CAS-object-barriers: combined problem of above



Shenandoah 2.0 Barriers: LRB

Load-reference-barriers:

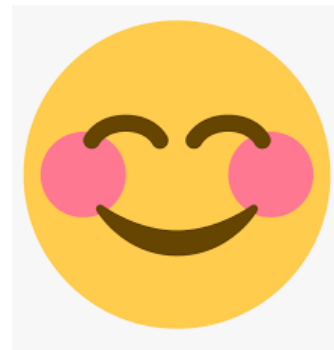
- LRB: Intercept loads of objects, establish canonical copy by CASing forwarding pointer
- CAS-object-barriers: still necessary but simpler



Shenandoah 2.0 Barriers: LRB

Load-reference-barriers:


- Much simpler barrier scheme
- Strong invariant
- Much less frequent (only object-loads!)
- Much simpler optimization story



Shenandoah 2.0 Barriers: LRB

Load-reference-barriers:

- Much simpler barrier scheme
 - Makes porting much easier
- Strong invariant
- Much less frequent (only object-loads)
 - Enables elimination of forwarding pointer
- Much simpler optimization story



Shenandoah Architectures


x86_64	✓	Primary target; continuously tested.
x86_32	✓	Secondary target; continuously tested.
AArch64	✓	Primary target; continuously tested.
ARM32	★	In development; help welcome.
PPC64	✗	Not supported; contributions welcome.
S390X	✗	Not supported; contributions welcome.
SPARC	✗	Not supported. No hardware to test on.
Others	✗	Please contact Shenandoah devs for guidance if you are willing to port Shenandoah to another platform.

Shenandoah OSes

Linux	✓	Primary target; continuously tested.
Windows	✓	Secondary target; continuously tested.
macOS	✓	Additional target; tested by community.
Solaris	✓	Additional target; tested by community.
Others	?	The porting should be trivial, please try and contact Shenandoah devs with your success and failure reports.

New Forwarding Pointer

Concurrent Compaction (Shenandoah 1.0):



New Forwarding Pointer

Shenandoah 1.0 – weak invariant:

- All reads **may** read from old copy (if no canonical copy has been established yet)
- All writes **must** write to new copy

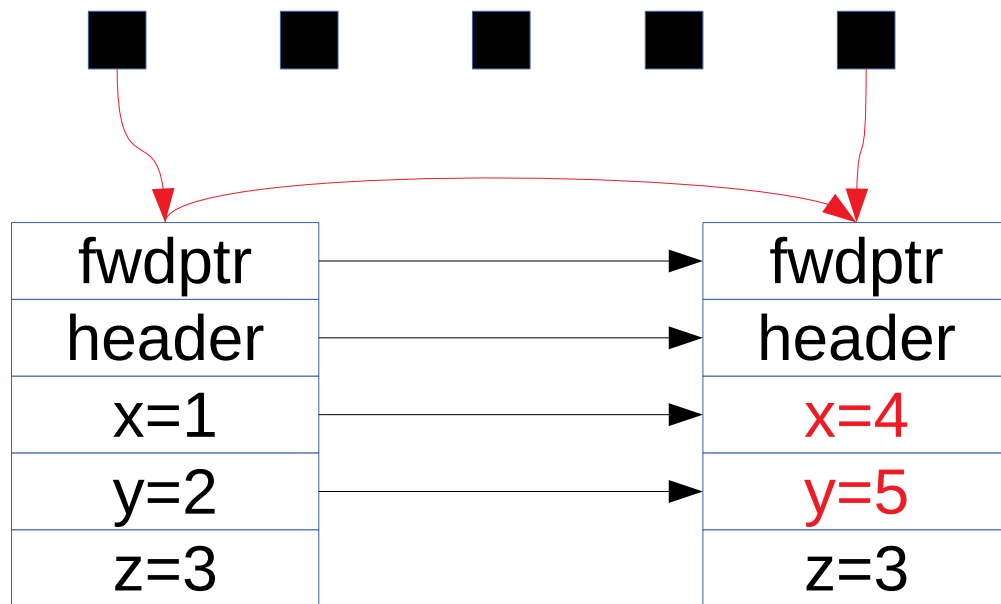
New Forwarding Pointer

Shenandoah 2.0 – strong invariant:

- All reads **must** read from new copy
- All writes **must** write to new copy
 - old copy is 100% unused
 - can use it to keep forwarding information


New Forwarding Pointer

Concurrent Compaction (Shenandoah 1.0):



New Forwarding Pointer

Concurrent Compaction (Shenandoah 2.0):



New Forwarding Pointer

```
T load_ref_barrier(T* addr) {
    T obj = *addr;
    if (in_cset(obj)) {
        T fwd = decode_fwdptr(obj);
        if (oop == fwd) {
            fwd = lrb_slowpath(oop);
        }
        return fwd;
    }
    return obj;
}
```

New Forwarding Pointer

```
// Shenandoah 1.0 impl
T decode_fwdptr(T obj) {
    return ((T*)obj)[fwdptr_offset];
}
```

```
// Shenandoah 2.0 impl
T decode_fwdptr(T obj) {
    intptr_t hdr = ((intptr_t*)obj)[header_offset];
    if ((hdr & 3) == 3) { // Forwarded
        return (T)(hdr & ~3);
    } else {
        Return obj;
    }
}
```

New Forwarding Pointer

- Reduces memory footprint per-object
- Down to 0.66x (best case), realistically 0.75x-0.95x
 - reduces allocation pressure
 - fewer GC cycles for same amount of allocs

Self-fixing Barriers

```
T load_ref_barrier(T* addr) {  
    T obj = *addr;  
    if (in_cset(obj)) {  
        T fwd = decode_fwdptr(obj);  
        if (oop == fwd) {  
            fwd = lrb_slowpath(oop);  
        }  
        return fwd;  
    }  
    return obj;  
}
```


Self-fixing Barriers

```
T load_ref_barrier(T* addr) {  
    T obj = *addr;  
    if (in_cset(obj)) {  
        T fwd = decode_fwdptr(obj);  
        if (oop == fwd) {  
            fwd = lrb_slowpath(oop);  
        }  
        return fwd;  
    }  
    return obj;  
}
```

Infrequent case: Object not
forwarded yet

Self-fixing Barriers

```
T load_ref_barrier(T* addr) {  
    T obj = *addr;  
    if (in_cset(obj)) {  
        T fwd = decode_fwdptr(obj);  
        if (oop == fwd) {  
            fwd = lrb_slowpath(oop);  
        }  
        return fwd;  
    }  
    return obj;  
}
```

Frequent case: Object forwarded,
but field not updated, yet

Self-fixing Barriers

- While we're there, couldn't we update the field?

Self-fixing Barriers

```
T load_ref_barrier(T* addr) {  
    T obj = *addr;  
    if (in_cset(obj)) {  
        T fwd = decode_fwdptr(obj);  
        if (oop == fwd) {  
            fwd = lrb_slowpath(oop);  
        } else {  
            cas(addr, fwd, obj);  
        }  
        return fwd;  
    }  
    return obj;  
}
```

Frequent case: Object forwarded,
field updated already

Availability

- Shenandoah 1.0 in JDK12
- Shenandoah 2.0 in JDK13
- 2.0 in shenandoah/jdk11 (Fedora, RHEL, ...)
- 2.0 in shenandoah/jdk8 (Fedora, RHEL, ...)
- Upstream 2.0 to jdk11u proper in progress
- Upstream 2.0 to jdk8u proper soon

Contacts

- Wiki:
<https://wiki.openjdk.java.net/display/shenandoah/Main>
- Mailing list:
<https://mail.openjdk.java.net/mailman/listinfo/shenandoah-dev>
- Twitter: @rkennke @shipilev

Questions?