



# Lilliput

## The Journey

Roman Kennke (@rkennke)

Principal Engineer  
Amazon




# What is Lilliput about?

Reduce memory footprint of Java applications


Reduce size of object headers from 12 to 4 bytes

Typically adds up to 20-30% heap reductions

# What is Lilliput about?



# What is Lilliput about?



# 2020

FOSDEM 2020

JDK 14 around the corner

Shenandoah 2.0

How hard can it be?

# 2021

First commit in March 2021

Lots of prototyping

Class-Pointer improvements

# 2022

April: First attempts to replace stack-locking

- > Finished in May 2023 (JDK 21)
- > Default since March 2024 (JDK 23)
- > Recursive locking support in April 2024 (JDK 23)
- > Object monitor table in August 2024

# 2023

March: Finished Lightweight locking (JDK 21)

May: First proposed version of JEP 450 implementation



# 2024

April: Recursive lightweight locking

August: Object monitor table

November: Integrated JEP 450 into JDK 24

# 2025

January:     Prototype Lilliput 2 aka 4-byte-headers

March:       EA of JDK 24

# What's in it?

## Current state

Header	
Klass*	Fields
Fields	
Fields	
Fields	
Fields	

Header	
Klass*	length
Elements	
Elements	
Elements	
Elements	

# What's in it?

## Lilliput 1

Header
Fields
Fields
Fields
Fields
Fields

Header	
length	Elms
Elements	
Elements	
Elements	
Elements	

# What's in it?

## Lilliput 2

Header	Fields
Fields	
Fields	
Fields	
Fields	

Header	length
Elements	
Elements	
Elements	
Elements	

# What's in it?

## Mark Word (normal):

64  
[.....HHHHHHHHHHHHHHHHHHHHHHHHHHHHHH.AAAA.TT]  
(Unused) (Hash Code) (GC Age)(Tag)

## Class Word (compressed):

[illegible]

## Insight:

- Most objects never get i-hashed
- Most objects never get locked



# Lilliput 1

Header (compact):

64	42	7	3	0
[CCCCCCCCCCCCCCCCCCCCCHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHV VVVVAAAASTT]				
(Compressed Class Pointer)		(Hash Code)	(GC Age)	^(Tag)
		(Self Forwarded Tag)		

- 22 Klass bits
- 31 Hash bits
- 4 Valhalla reserved bits
- 4 GC age bits
- 2+1 Tag bits



# Lilliput 2 – The Plan

Header (Lilliput 2):

```
32          11    7    3    0
[CCCCCCCCCCCCCCCCCHVVVAAAATT]
  (Class Pointer)    ^ (Age) ^ (Tag)
                    (Hash-Code) (Self Forwarded Tag)
```

- 19 Klass bits
- 2 Hash control bits
- 4 Valhalla reserved bits
- 4 GC age bits
- 2+1 Tag bits

# Compact Identity Hash-Code

- Recap:

- $x == y \Rightarrow \text{System.identityHashCode}(x) == \text{System.identityHashCode}(y)$
- '==' is object identity, not equality
- Inverse not necessarily true
- Id-hash must remain stable throughout object lifetime (idempotence)
- Shoot for great entropy
- `System.identityHashCode()` is default impl for `Object.hashCode()`

# Compact Identity Hash-Code

- Current situation:
  - Allocate 31 bits slot for every object
  - First call to `System.identityHashCode()` installs good hash (RNG)
  - Subsequent calls read and return that hash
  - Only small fraction of objects ever use it

# Compact Identity Hash-Code

- Naïve Idea:
  - Calculate I-hash based on object address
- Problem:
  - GCs can move objects around
  - -> not idempotent

# Compact Identity Hash-Code

- Idea:
  - Allocate slot for id-hash on-demand
- Problem:
  - Heap is filled contiguously
  - There is no space available



# Compact Identity Hash-Code

- 00: Not-hashed/not-expanded (initial state)
- First call to `System.identityHashCode()` on an object:
  - Transitions state to 01: hashed/not-expanded
  - Calculates hash-code based on object address

# Compact Identity Hash-Code

- 01: Hashed/Not-expanded
- Subsequent calls to `System.identityHashCode()`:
  - Re-calculate hash-code based on object address



# Compact Identity Hash-Code

- First GC after object transitioned to 01: hashed/not-expanded:
  - Calculate hash-code based on original object address
  - Copy object, maybe extend to make space for i-hash-slot
    - No need to expand if alignment gap is available
  - Store i-hash in i-hash-slot
  - Transition to state 11: hashed/expanded

# Compact Identity Hash-Code

- 11: Hashed/Expanded
- Subsequent calls to `System.identityHashCode()`:
  - Read i-hash from i-hash-slot
- I-hash slot precomputed by field layouter
  - Can often use gaps

# Compact Identity Hash-Code

Header (Lilliput 2):

```
32          11    7    3    0
[CCCCCCCCCCCCCCCCCHVVVAAAATT]
  (Class Pointer)    ^ (Age) ^ (Tag)
                    (Hash-Code) (Self Forwarded Tag)
```

- Requires only 2 bits per object
- Allocate space for i-hash only when needed
- Majority of objects never need i-hash
- ~50% of objects have gaps available

# GC Forwarding

- Current situation:

## Mark Word (normal):

64 39 8 3 0  
[.....HHHHHHHHHHHHHHHHHHHHHHHHHHHHHH.AAAA.TT]  
(Unused) (Hash Code) (GC Age) (Tag)


## Mark Word (GC forwarded):

[illegible]

# GC Forwarding

- 2 cases:
  - Copying compaction (non-destructive)
  - Sliding compaction (destructive)

# GC Forwarding – Copying Compaction



# GC Forwarding – Copying Collection


- 3 phases:
  1. Marking (find live objects)
  2. Copy objects to empty space (write forwarding pointer into old object)
  3. Update all references

# GC forwarding - Copying Compaction

- Easy: keep using first word as forwarding pointer
- Overwrites first field(s) of object (ok, we have a copy)
- Overwrites length of arrays (ok, but we need to be careful)



# GC Forwarding - Sliding Compaction



# GC Forwarding - Sliding Compaction

- 4 phases:
  1. Marking (find live objects)
  2. Calculate forwarding addresses (write them into object headers)
  3. Update all references
  4. Copy objects to new locations

# GC Forwarding - Sliding Compaction

- Problem:
  - Forwarding overwrites Klass\* (and array-length, hash-bits)
  - ... but we don't have a copy
  - ... and we only have 9 bits

Header (Lilliput 2):

```
32                                11   7   3   0
[CCCCCCCCCCCCCCCCCCCCCHHHHHHHHHFTT]
      (Class Pointer)      ^ (Age) ^ (Tag)
                          (Hash-Code) (Self
```

Forwarded Tag)

# GC Forwarding – Sliding Compaction

- Idea:
  - 9 bits are enough to address  $2^9 = 512$  words range
  - Divide heap into 512 words (=4KB) blocks
  - Use side-table with one entry per block (=  $1/512^{\text{th}}$  of heap size)
  - Table stores forwarding base for each block
    - It's the forwarding address of the first live object in the block
  - Use 9 header bits to find forwarding address of object

$$\text{fwd\_addr} = \text{fwd\_base}[\text{block}] + \text{fwd\_bits} \ll 3$$

# GC Forwarding – Sliding Forwarding

- Fixed-overhead side table ( $1/512^{\text{th}}$  of heap size)
- No performance loss
- Potential to reduce to  $1/2048^{\text{th}}$  of heap size, by using tag bits
- Eliminates 8TB limit of Lilliput 1

# Putting it all together

- Tiny Class-Pointers
  - Allow 19 bit class-pointers
- Compact Identity Hash-Code
  - Need only 2 header bits for i-hash
- Improved full-GC forwarding table
- 4 Valhalla bits

- Header (Lilliput 2):

32 11 7 3 0  
[CCCCCCCCCCCCCCCCCHHVVVVAAAATT]  
(Class Pointer) ^ (Age) ^ (Tag)  
(Hash-Code) (Self Forwarded Tag)

# Putting it all together

- `-XX:+UseCompactObjectHeaders`